

Deep Learning + Neural Networks + Machine Learning + TensorFlow + Keras +

Deep Learning with Python: An Introduction to TensorFlow and Keras



Muhammad Taha

Follow

3 min read · Jul 12, 2024



41



In the rapidly evolving field of artificial intelligence, deep learning stands out as one of the most transformative technologies. By mimicking the human brain, deep learning models can recognize patterns in vast amounts of data, enabling groundbreaking advancements in areas like image recognition, natural language processing, and beyond. In this article, we'll explore the basics of deep learning, get you set up with TensorFlow and

Keras, and guide you through building and training your first neural network.

Basics of Deep Learning and Neural Networks

What is Deep Learning?

Deep learning is a subset of machine learning that utilizes neural networks with many layers (hence the term “deep”). These layers enable the model to learn from data in a hierarchical fashion, extracting increasingly complex features at each layer. This approach is particularly powerful for tasks involving large datasets and complex patterns.

Neural Networks: The Building Blocks

A neural network is composed of layers of interconnected nodes, or neurons. Each neuron takes an input, applies a weighted sum and an activation function, and passes the output to the next layer. The network learns by adjusting these weights based on the error in its predictions, a process called backpropagation.

Get Muhammad Taha's stories in your inbox

Join Medium for free to get updates from this writer.

Enter your email

Subscribe

☐ Remember me for faster sign in

Here's a simple diagram of a neural network:




Image by [GeeksforGeeks](#)

Key Concepts:

- **Layers:** Collections of neurons. Common types include dense (fully connected) layers, convolutional layers (for image data), and recurrent layers (for sequential data).

- **Activation Functions:** Functions that introduce non-linearity into the model, enabling it to learn complex patterns. Examples include ReLU, sigmoid, and tanh.
- **Loss Function:** Measures the error between the model's prediction and the actual target. Common loss functions include mean squared error (MSE) for regression and categorical cross-entropy for classification.
- **Optimizer:** Algorithm that adjusts the model's weights to minimize the loss function. Popular optimizers include SGD, Adam, and RMSprop.

Setting Up TensorFlow and Keras

TensorFlow and Keras are two popular frameworks for building and training deep learning models. TensorFlow is an open-source platform developed by Google, while Keras is a high-level neural networks API, written in Python and capable of running on top of TensorFlow.

Installation

To get started, you'll need to install TensorFlow and Keras. You can do this easily using pip:

```
pip install tensorflow keras
```

Verifying the Installation

After installation, verify that TensorFlow and Keras are correctly installed by running the following Python code:

```
import tensorflow as tf
from tensorflow import keras

print(tf.__version__)
print(keras.__version__)
```

If you see version numbers for both TensorFlow and Keras, you're good to go!

Building and Training Your First Neural Network

Let's build a simple neural network to classify handwritten digits from the famous MNIST dataset. This dataset contains 60,000 training images and

10,000 test images of handwritten digits (0–9), each represented as a 28x28 pixel grid.

Step 1: Import Libraries and Load Data

```
import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt

# Load the MNIST dataset
(train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.mnist.load_data()

# Normalize the images to [0, 1] range
train_images = train_images / 255.0
test_images = test_images / 255.0
```

Step 2: Build the Model

We'll create a simple sequential model with two dense (fully connected) layers. The first layer will have 128 neurons with a ReLU activation function, and the second layer will be the output layer with 10 neurons (one for each digit) and a softmax activation function.

```
model = models.Sequential([  
    layers.Flatten(input_shape=(28, 28)), # Flatten the 28x28 images into 1D ve  
    layers.Dense(128, activation='relu'),  
    layers.Dense(10, activation='softmax')  
])
```

Step 3: Compile the Model

Next, we need to compile the model by specifying the optimizer, loss function, and metrics to monitor.

```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

Step 4: Train the Model

We can now train the model using the training data. We'll also validate the model on the test data to see how well it generalizes.


```
model.fit(train_images, train_labels, epochs=5, validation_data=(test_images, te
```

Step 5: Evaluate the Model

Finally, let's evaluate the model's performance on the test dataset.

```
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f'Test accuracy: {test_acc}')
```

Conclusion

Congratulations! You've just built and trained your first neural network using TensorFlow and Keras. This is just the beginning — deep learning has a vast array of techniques and architectures to explore. With TensorFlow and Keras, you have powerful tools at your disposal to dive deeper into the world of artificial intelligence.

Feel free to share your experiences and insights in the comments below.
Happy coding!

Deep Learning

Neural Networks

Machine Learning

TensorFlow

Keras



Written by Muhammad Taha

4 followers · 1 following

Follow

MERN Fellow @ Bytewise Limited | Student of BS Software Engineering @NUML |
Artificial Intelligence | Machine Learning

No responses yet



Write a response

What are your thoughts?

More from Muhammad Taha




 Muhammad Taha · Jul 11, 2024

Building Your First Machine Learning Model with Scikit-learn

In this guide, we'll walk through the process of building your first machine learning model...

 2  1




 Muhammad Taha · Jul 18, 2024

Ethics and Bias in AI: Ensuring Fairness in Machine Learning...

Artificial Intelligence (AI) has revolutionized numerous fields, offering unprecedented...






 Muhammad Taha · Jul 17, 2024

Deploying Machine Learning Models with Flask and Docker

Deploying machine learning models can often seem daunting, but with the right tools...



 Muhammad Taha · Jul 16, 2024


Hyperparameter Tuning in Machine Learning: Best Practices

In the world of machine learning, the success of a model often hinges on its...



See all from Muhammad Taha

Recommended from Medium



 Ritik Jain · Jan 20

The Complete Guide to Fine-Tuning Embedding Models: From Theory...

A comprehensive guide to understanding, preparing data, loss function, eval metrics fo...

★ 🖱️ 128 💬 2 ↺ 1 

 In Write A Catal... by Dr. Patricia Schmidt... · Jan 15

As a Neuroscientist, I Quit These 5 Morning Habits That Destroy You...


Most people do #1 within 10 minutes of waking (and it sabotages your entire day)

★ 🖱️ 60K 💬 1376 ↺ 582 



 In Data Science Collective by Marina Wyss · Feb 23

Should You Still Learn to Code in 2026?




 In Towards AI by Shreyas Naphad · Mar 29

If You Understand These 5 AI Terms, You're Ahead of 90% of...

The answer isn't as obvious as I used to believe.

★ 🖐️ 5.5K 💬 282 ↻ 222 💬



 Leo Godin · Mar 2

Claude Code is Great

You Just Need to Learn How to Use It

★ 🖐️ 6.1K 💬 186 ↻ 141 💬

Master the core ideas behind AI without getting lost

★ 🖐️ 21K 💬 526 ↻ 175 💬



 In UX Planet by Nick Babich 🇺🇸 · Apr 6

MCP is Dead

Why you should avoid using MCP in Claude Code and what to use instead

★ 🖐️ 5K 💬 267 ↻ 172 💬

See more recommendations

[Help](#) [Status](#) [About](#) [Careers](#) [Press](#) [Blog](#) [Store](#) [Privacy](#) [Rules](#) [Terms](#) [Text to speech](#)